


**Exercise 24.** For each of the histories shown in Figs. 3.13 and 3.14 are they quiescently consistent? Sequentially consistent? Linearizable? Justify your answer.



**Figure 3.13** First history for Exercise 24.

The following history is sequential consistent since a possible execution order could be: write(1), read(1), write(2), read(2). Additionally, the history is quiescently consistent the method calls to appear in one-at-a-time sequential order and there are no pending invocations as seen in Figure 3.13. Finally, the history is linearizable since we can have a history write(1), read(1), write(2), read(2) such that the writes complete before the reads correctly return 1 then 2.



**Figure 3.14** Second history for Exercise 24.

The following history is sequential consistent since a possible execution order could be: write(2), write(1), read(1), read(1). Additionally, the history is quiescently consistent since the method calls appear in a one-at-a-time sequential order and there are no pending invocations as seen in Figure 3.14. Finally, the history is linearizable since we can have a history write(2), write(1), read(1), read(1) such that the writes complete before the reads and the reads correctly return 1 twice.

**Exercise 25.** If we drop condition L2 from the linearizability definition, is the resulting property the same as sequential consistency? Explain.

Condition L2 states that if a method call  $m_0$  precedes method call  $m_1$  in  $H$ , then the same is true in  $S$  where  $H$  is a history and  $S$  is a legal sequential history that is equivalent to the complete extension of  $H$ . Condition L2 is the quiescent property for linearization, which makes linearizability compositional. Condition L1 states a previous method call must have taken effect before a later method call and Sequential Consistency requires that method calls act as if they occurred in a sequential order consistent with program order. Condition L1 is essential the sequential consistency property for linearizability. Therefore, removing condition L2 from linearizability results in a property equivalent to sequential consistency.

**Exercise 27.** The `AtomicInteger` class (in the `java.util.concurrent.atomic` package) is a container for an integer value. One of its methods is

```
boolean compareAndSet(int expect, int update).
```

This method compares the object's current value to `expect`. If the values are equal, then it atomically replaces the object's value with `update` and returns *true*. Otherwise, it leaves the object's value unchanged, and returns *false*. This class also provides

```
int get()
```

which returns the object's actual value.

Consider the FIFO queue implementation shown in Fig. 3.15. It stores its items in an array `items`, which, for simplicity, we will assume has unbounded size. It has two `AtomicInteger` fields: `tail` is the index of the next slot from which to remove an item, and `head` is the index of the next slot in which to place an item. Give an example showing that this implementation is *not* linearizable.

```
1 class IQueue<T> {
2   AtomicInteger head = new AtomicInteger(0);
3   AtomicInteger tail = new AtomicInteger(0);
4   T[] items = (T[]) new Object[Integer.MAX_VALUE];
5   public void enq(T x) {
6     int slot;
7     do {
8       slot = tail.get();
9     } while (! tail.compareAndSet(slot, slot+1));
10    items[slot] = x;
11  }
12  public T deq() throws EmptyException {
13    T value;
14    int slot;
15    do {
16      slot = head.get();
17      value = items[slot];
18      if (value == null)
19        throw new EmptyException();
20    } while (! head.compareAndSet(slot, slot+1));
21    return value;
22  }
23 }
```

Figure 3.15 IQueue implementation.

Given threads 1 and 2, thread 1 `enq(A)`, but stops executing before setting `item[0]` and never finishes. Thread 2 calls `enqueue(B)` and assigns `item[1]` to B since thread 1 was supposed to assign `item[0]` and the `get()` method returned 1. Next, Thread 2 calls `dequeue()` and attempts to remove A from `item[0]`, but throws an empty exception because thread 1 never completed executing and did not set `item[0]` to A.

The example proves the implementation is not linearizable, since `item[0]` should contain a value, A, instead of throwing an empty exception. The implementation violates the condition that if one method call precedes another, then the earlier call must have taken effect before the later call.

**Exercise 32.** This exercise examines a queue implementation (Fig. 3.17) whose `enq()` method does not have a linearization point.

The queue stores its items in an `items` array, which for simplicity we will assume is unbounded. The `tail` field is an `AtomicInteger`, initially zero. The `enq()` method reserves a slot by incrementing `tail`, and then stores the item at that location. Note that these two steps are not atomic: there is an interval after `tail` has been incremented but before the item has been stored in the array.

The `deq()` method reads the value of `tail`, and then traverses the array in ascending order from slot zero to the tail. For each slot, it swaps `null` with the current contents, returning the first non-`null` item it finds. If all slots are `null`, the procedure is restarted.

Give an example execution showing that the linearization point for `enq()` cannot occur at Line 15.

Hint: give an execution where two `enq()` calls are not linearized in the order they execute Line 15.

Proof by Counterexample:

Given threads A, B, C, thread A enqueues 'a' into the queue concurrently with thread B enqueueing 'b', then thread C dequeues.

- 1) Thread A executes `tail.getAndIncrement()` setting 'i' to 0 and `tail` to 1.
- 2) Thread B executes `tail.getAndIncrement()` setting 'i' to 1 and `tail` to 2.
- 3) Thread B sets `items[1]` to 'b'.
- 4) Thread C dequeues at `items[0]` returning null.
- 5) Thread C moves to `items[1]` and dequeues returning 'b'.
- 6) Thread A sets `items[0]` to 'a'.
- 7) Thread C dequeues at `items[0]` returning 'a'.

The possible history demonstrates a counterexample where 'b' is returned before 'a' when thread A executes line 15 before thread B, therefore line 15 is not a linearization point since it is not visible to other threads.

```

1 public class HWQueue<T> {
2     AtomicReference<T>[] items;
3     AtomicInteger tail;
4     static final int CAPACITY = 1024;
5
6     public HWQueue() {
7         items = (AtomicReference<T>[]) Array.newInstance(AtomicReference.class,
8             CAPACITY);
9         for (int i = 0; i < items.length; i++) {
10             items[i] = new AtomicReference<T>(null);
11         }
12         tail = new AtomicInteger(0);
13     }
14     public void enq(T x) {
15         int i = tail.getAndIncrement();
16         items[i].set(x);
17     }
18     public T deq() {
19         while (true) {
20             int range = tail.get();
21             for (int i = 0; i < range; i++) {
22                 T value = items[i].getAndSet(null);
23                 if (value != null) {
24                     return value;
25                 }
26             }
27         }
28     }
29 }

```

Figure 3.17 Herlihy/Wing queue.

Give another example execution showing that the linearization point for `enq()` cannot occur at Line 16.

Since these are the only two memory accesses in `enq()`, we must conclude that `enq()` has no single linearization point. Does this mean `enq()` is not linearizable?

Given threads A, B, C, thread A executes concurrently with thread B, then thread C dequeues.

- 1) Thread A enqueues, executes `tail.getAndIncrement()`, so thread A gets `items[0]`.
- 2) Thread B enqueues after thread A calling `tail.getAndIncrement()` and setting `items[1]` to 'b' before thread A sets `items[0]` to 'a'.
- 3) Thread C dequeues at `items[0]` returning 'a'.
- 4) Thread C dequeues again at `items[1]` returning 'b'.

Line 16 is not a linearization point since thread B called `items[i].set(x)`, line 16, before thread A even though thread A was first in the queue. Thread A called `getAndIncrement` before thread B, so it had `items[0]` even though it had not assigned 'a' to `items[0]` before thread B assigned `items[1]` to 'b'.

The `enq()` method can be linearizable, but doesn't have single linearization point in this case.